


Pipelines
Tabular Data
Embeddings

Document Similarity




Why might you consider text messages similar?

We denote by $\vec{V}(d)$ the vector derived from document (message) d , with one component in the vector for each dictionary term. The set of documents in a collection then may be viewed as a set of vectors in a vector space, in which there is one axis for each term.

The similarity between two documents d_1 and d_2 can be found by the cosine similarity of their vector representations $\vec{V}(d_1)$ and $\vec{V}(d_2)$:

$$S_{12} = \frac{\vec{V}(d_1) \cdot \vec{V}(d_2)}{|\vec{V}(d_1)| \times |\vec{V}(d_2)|}$$

A collection then is a **term-document** matrix. For example, terms in various "period" novels.



term	SaS	PaP	WH	term	SaS	PaP	WH
affection	115	58	20	affection	0.996	0.993	0.847
jealous	10	7	11	jealous	0.087	0.120	0.466
gossip	2	0	6	gossip	0.017	0	0.254

Consider the query $q = \text{jealous gossip}$. This query turns into the unit vector $\vec{V}(q) = (0, 0.707, 0.707)$.

Cosine similarity is the dot product of unit vectors.

Wuthering Heights is the top-scoring document for this query with a score of 0.509.

Embeddings: extend the idea

- Observe a bunch of people
- **Infer** Personality traits from them
- Vector of traits is called an **Embedding**
- Who is more similar? Jay and who?
- Use Cosine Similarity of the vectors

	Trait #1	Trait #2	Trait #3	Trait #4	Trait #5
Jay	-0.4	0.8	0.5	-0.2	0.3
Person #1	-0.3	0.2	0.3	-0.4	0.9
Person #2	-0.5	-0.4	-0.2	0.7	-0.1

$\text{cosine_similarity}(\text{Jay}, \text{Person \#1}) = 0.66$ ✓

$\text{cosine_similarity}(\text{Jay}, \text{Person \#2}) = -0.37$

Categorical Data

Example:

Rossmann Kaggle Competition. Rossmann is a 3000 store European Drug Store Chain. The idea is to predict sales 6 weeks in advance.

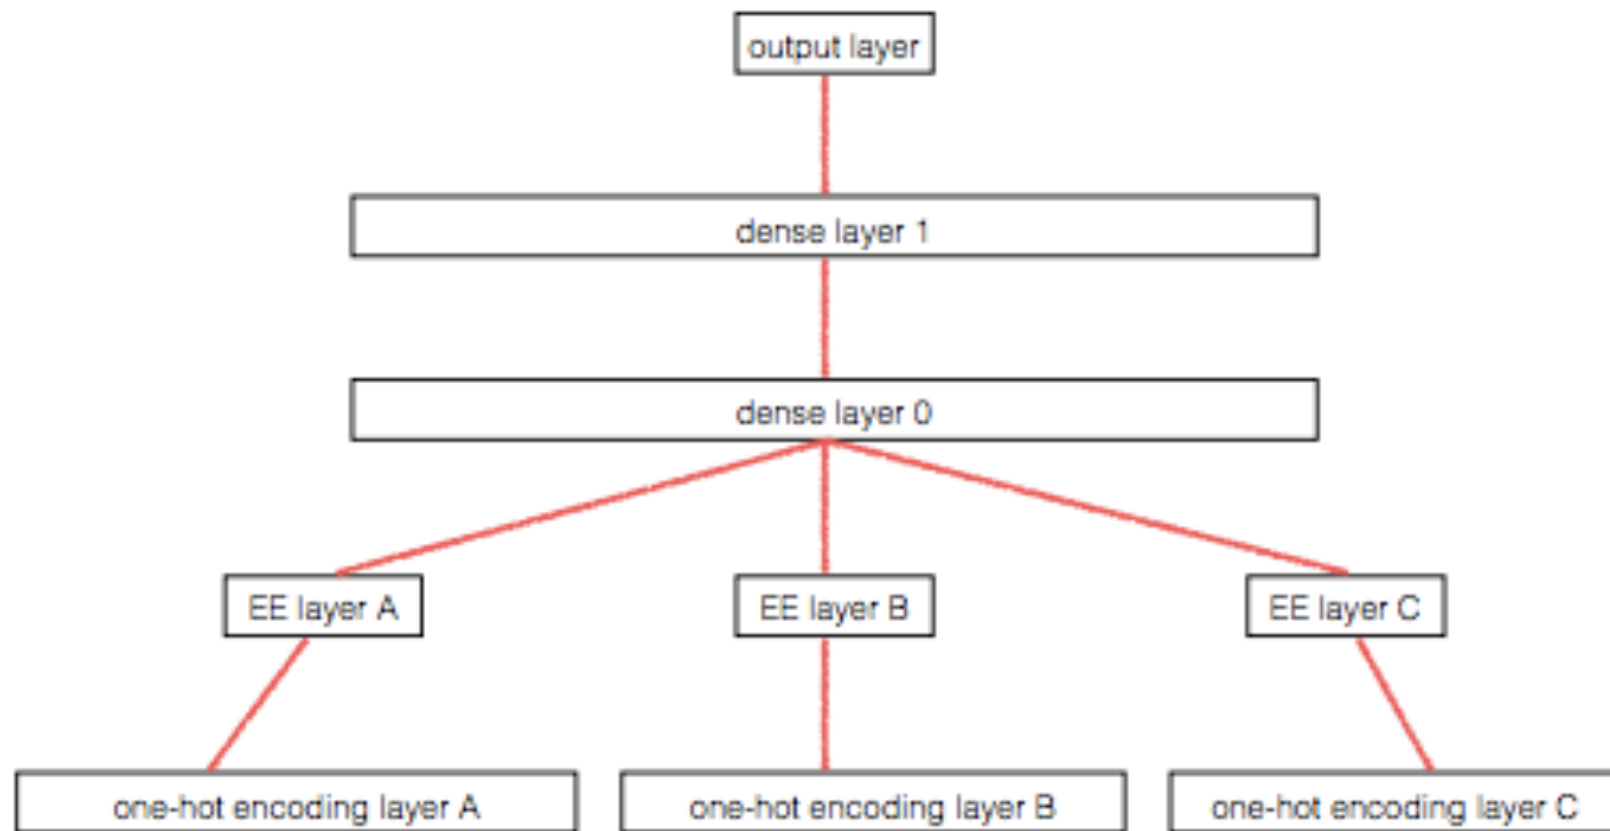
Consider `store_id` as an example. This is a **categorical** predictor, i.e. values come from a finite set.

We usually **one-hot encode** this: a single store is a length 3000 bit-vector with one bit flipped on.

What is the problem with this?

- The 3000 stores have commonalities, but the one-hot encoding does not represent this
- Indeed the dot-product (cosine similarity) of any-2 1-hot bitmaps must be 0
- Would be useful to learn a lower-dimensional **embedding** for the purpose of sales prediction.
- These store "personalities" could then be used in other models (different from the model used to learn the embedding) for sales prediction
- The embedding can be also used for other **tasks**, such as employee turnover prediction

Training an Embedding



- Normally you would do a linear or MLP regression with sales as the target, and both continuous and categorical features
- The game is to replace the 1-hot encoded categorical features by "lower-width" embedding features, for *each* categorical predictor
- This is equivalent to considering a neural network with the output of an additional **Embedding Layer** concatenated in
- The Embedding layer is simply a linear regression


Training an embedding (contd)

A 1-hot vector for a categorical variable v with cardinality $N(v)$ can be written using the Kronecker Delta symbol as

$$v_k = \delta_{jk}, j \in \{1..N(v)\}$$

Then an embedding of width (dimension) L is just a $N(v) \times L$ matrix of weights W_{ij} such that multiplying the k th 1-hot vector by this weight matrix picks out the k th row of weights (see right)

But how do we find these weights? We fit for them with the rest of the weights in the MLP!




```

def build_keras_model():
    input_cat = []
    output_embeddings = []
    for k in cat_vars+nacols_cat: #categoricals plus NA booleans
        input_1d = Input(shape=(1,))
        output_1d = Embedding(input_cardinality[k], embedding_cardinality[k], name='{}_embedding'.format(k))(input_1d)
        output = Reshape(target_shape=(embedding_cardinality[k],))(output_1d)
        input_cat.append(input_1d)
        output_embeddings.append(output)

    main_input = Input(shape=(len(cont_vars),), name='main_input')
    output_model = Concatenate()([main_input, *output_embeddings])
    output_model = Dense(1000, kernel_initializer="uniform")(output_model)
    output_model = Activation('relu')(output_model)
    output_model = Dense(500, kernel_initializer="uniform")(output_model)
    output_model = Activation('relu')(output_model)
    output_model = Dense(1)(output_model)

    kmodel = KerasModel(
        inputs=[*input_cat, main_input],
        outputs=output_model
    )
    kmodel.compile(loss='mean_squared_error', optimizer='adam')
    return kmodel

def fitmodel(kmodel, Xtr, ytr, Xval, yval, epochs, bs):
    h = kmodel.fit(Xtr, ytr, validation_data=(Xval, yval),
                    epochs=epochs, batch_size=bs)

    return h

```

Another Example: Recommendation Systems


We want to make a recommendation (on, say, a 1-5 scale) for an item m by a user u .

We can write this as: $Y_{um} = Y_{um}^{baseline}$

where $Y_{um}^{baseline} = \mu + \bar{\theta} \cdot I_u + \bar{\gamma} \cdot I_m$

where the unknown parameters θ_u and γ_m indicate the deviations, or biases, of user u and item m , respectively, from some intercept parameter μ .

$$Y_{35}^{baseline} = \mu + \bar{\theta} \cdot I_3 + \bar{\gamma} \cdot I_5 = \mu + [I_3, I_5] \cdot [\bar{\theta}, \bar{\gamma}]$$



Remember this is a sparse problem. Most users have not rated most items. So we will want to regularize. Thus we want to minimize the loss:


$$\sum_{u,m} \left(Y_{um} - \mu - \bar{\theta} \cdot I_u - \bar{\gamma} \cdot I_m \right)^2 + \alpha \left(\theta_u^2 + \gamma_m^2 \right)$$

This is a Ridge Regression, or SGD with weight decay.

The baseline is not enough, though. Lets add a term r_{um} , the residual left after the biases are taken into account:

$$Y_{um} = Y_{um}^{baseline} + r_{um}.$$

L=2 latent factors



Modeling the Residual

Associate with each item a vector \bar{q}_m of length L . Associate with each user a vector \bar{p}_u of length L .

\bar{q}_m measures the extent to which an item(restaurant) possesses L latent factors: Low price, spicyness, etc. \bar{p}_u captures the extent of interest the user has in restaurants(items) that are high on the corresponding factors (a user who likes spicy food).

Then we model the residuals as:

$$r_{um} = \bar{q}_m^T \cdot \bar{p}_u,$$

the user's overall interest in the item's characteristics.

Embeddings Again

So, we want: $Y_{um} = \mu + \bar{\theta} \cdot I_u + \bar{\gamma} \cdot I_m + \bar{q}_m^T \cdot \bar{p}_u$

To solve this we need to simply minimize the risk of the entire regression, ie

$$\sum_{u,m} \left(Y_{um} - \mu - \bar{\theta} \cdot I_u - \bar{\gamma} \cdot I_m - \bar{q}_m^T \cdot \bar{p}_u \right)^2 + \alpha \left(\theta_u^2 + \gamma_m^2 + \|\bar{q}_m\|^2 + \|\bar{p}_u\|^2 \right)$$

We have seen the \bar{q}_m^T and \bar{p}_u before! These are simply L dimensional item and user specific **embeddings!!** And we'll train the entire model using SGD and weight decay. (The biases can be thought as 1-D embeddings!)

Embedding is just a linear regression

So why are we giving it another name?

- it is usually to a lower dimensional space
- traditionally we have done linear dimensional reduction through PCA or SVD and truncation, but sparsity can throw a spanner into the works
- we train the weights of the embedding regression using SGD, along with the weights of the downstream task (here fitting the rating)
- the embedding can be used for alternate tasks, such as finding the similarity of users.

See how **Spotify** does all this..

```
def embedding_input(emb_name, n_items, n_fact=20, l2regularizer=1e-4):  
    inp = Input(shape=(1,), dtype='int64', name=emb_name)  
    return inp, Embedding(n_items, n_fact, input_length=1, embeddings_regularizer=l2(l2regularizer))(inp)
```

```
usr_inp, usr_emb = embedding_input('user_in', n_users, n_fact=50, l2regularizer=1e-4)  
mov_inp, mov_emb = embedding_input('movie_in', n_movies, n_fact=50, l2regularizer=1e-4)
```

```
def create_bias(inp, n_items):  
    x = Embedding(n_items, 1, input_length=1)(inp)  
    return Flatten()(x)
```

```
usr_bias = create_bias(usr_inp, n_users)  
mov_bias = create_bias(mov_inp, n_movies)
```

```
def build_dp_bias_recommender(u_in, m_in, u_emb, m_emb, u_bias, m_bias):  
    x = dot([u_emb, m_emb], axes=(2,2))  
    x = Flatten()(x)  
    x = add([x, u_bias])  
    x = add([x, m_bias])  
    bias_model = Model([u_in, m_in], x)  
    bias_model.compile(Adam(0.001), loss='mse')  
    return bias_model
```

```
bias_model = build_dp_bias_recommender(usr_inp, mov_inp, usr_emb, mov_emb, usr_bias, mov_bias)
```




Word Embeddings

- The Vocabulary V of a corpus (large swath of text) can have 10,000 and maybe more words
- a 1-hot encoding is huge, moreover, similarities between words cannot be established
- we map words to a smaller dimensional latent space of size L by considering some downstream task to train on
- we hope that the embeddings learnt are useful for other tasks.

Obligatory example


See man->boy as woman->girl, similarities of king and queen, for eg. These are lower dimensional GloVe embedding vectors

king - man + woman ≈ queen



How do we train word embeddings?

We need to choose a downstream task. We could choose **Language Modeling**: predict the next word. We'll start with random "weights" for the embeddings and other parameters and run SGD. A trained model+embeddings would look like this:



How do we set up a training set?

Thou shalt not make a machine in the likeness of a human mind

Sliding window across running text

thou	shalt	not	make	a	machine	in	the	...
thou	shalt	not	make	a	machine	in	the	
thou	shalt	not	make	a	machine	in	the	
thou	shalt	not	make	a	machine	in	the	
thou	shalt	not	make	a	machine	in	the	

Dataset

input 1	input 2	output
thou	shalt	not
shalt	not	make
not	make	a
make	a	machine
a	machine	in

Why not look both ways? This leads to the Skip-Gram and CBOW architectures..

SKIP-GRAM: Predict Surrounding Words

Choose a window size (here 4) and construct a dataset by sliding a window across.

Thou shalt not make a machine in the likeness of a human mind

thou	shalt	not	make	a	machine	in	the	...
------	-------	-----	------	---	---------	----	-----	-----

input word	target word
not	thou
not	shalt
not	make
not	a

Thou shalt not make a machine in the likeness of a human mind

thou	shalt	not	make	a	machine	in	the	...
------	-------	-----	------	---	---------	----	-----	-----

thou	shalt	not	make	a	machine	in	the	...
------	-------	-----	------	---	---------	----	-----	-----

input word	target word
not	thou
not	shalt
not	make
not	a
make	shalt
make	not
make	a
make	machine

Usage of word2vec

- the pre-trained word2vec and other embeddings (such as GloVe) are used everywhere in NLP today
- the ideas have been used elsewhere as well. **AirBnB** and **Anghami** model sequences of listings and songs using word2vec like techniques
- **Alibaba** and **Facebook** use word2vec and graph embeddings for recommendations and social network analysis.